



-- Theory: attention scores only --●-- 4:1 model vs. baseline (measured)


Peak RAM



Step RAM



Time per step



Negative = the 4:1 model needs less. Measured: mean $\pm$ std of the per-repetition difference. Theory: total size of all attention-score tensors (the part sliding-window attention changes). Values whose baseline is below 0.05 (GB or s) are left out as too small to compare (step RAM at L=128).